



# Entity tagging of seminar announcements

Csongor Barabasi – 1636980

## 1. Outlining the Task

The task was, given 300 tagged emails, to build a tagger which will be able to tag 300 untagged emails accurately, in the same way the tagged emails were tagged. The corpus was tagged specifically with stime, etime, sentence, paragraph, location and speaker tags.

## 2. Data Pre-Processing

The first step was reading the corpus. This was done using *WordListCorpusReader* from *NLTK*. Then, using the *raw* property of the corpus reader, we could retrieve all the emails, concatenated. Before splitting it into email chunks, a cleaning process was done, where sequences of two or more characters like '=', '~', '%', '-', '\_', '\*', '!' were removed. Next step involved splitting this long text of 300 emails into emails chunks, and then split up each email into a header and a body part, which was then stored in an *Email* object. Finding each email was done with the regular expression

`<.*? @.*?\nType >`

split the text according to each email's header. Then, by splitting each email according to the word 'Abstract', I was able to split each email into a header and body part. The result of the pre-processing part was a list of *Email* objects.

## 3. Sentence and Paragraph Tagging

In the given corpus paragraphs can be identified as being separated by two newline characters. Hence, I could obtain a list of paragraphs from each email body by splitting the text as:

`text.split('\n\n')`

Out of those I kept only those paragraphs which do not start with a white space and the first character is alphanumeric. Evaluated on the test data, I was able to achieve a precision of **0.5**, recall of **0.57** leading to an F1-Score of **0.533**.

Sentence tagging was as simple as paragraph tagging. Using *NLTK*'s built-in *sent\_tokenize* function, which uses an instance of *PunktSentenceTokenizer*, I was able to split the text into sentences, and from those choose the ones which does not start with a white space and the first character is alphanumeric. After evaluation on the test data, this approach achieved a precision of **0.616**, recall of **0.104** and F1-Score of **0.178**.

In the future improvements could be made to both approaches by adding more features based on which to decide whether a part of text is a sentence or paragraph, mostly focusing on increasing the recall of both algorithms.

## 4. Time Tagging

First, I extract the line from the header which gives information about time using the regex

*Time.\*\n*

Then we extract all the existing timestamps from this line using the regular expression

*/[0 - 9]+: [0 - 9][0 - 9]||[0 - 9]+: [0 - 9][0 - 9] + [APap]\.?[mM]*

Next, I checked whether there were one or two timestamps found. If there was only one, it was tagged as *stime*, if there were two then the first was tagged *stime* and the second as *etime*. Next, I scanned through the body of the email and looked for all the occurrences of the times found in the header, and tagged them in the body accordingly (e.g. if there was one timestamp in the header, that was tagged with *stime*, so all occurrences of that timestamp in the body were tagged as *stime*).

Evaluating on the test data this approach performed really well, achieving a precision of **0.97**, recall of **0.48**, F1-Score of **0.646**.

## 5. Speaker Tagging

Based on the training data I constructed a database, which contains all the speakers which were tagged in the training data. As a first step I check whether there exists a line starting with the word 'Who:' in the header. In case it exists, I loop through my speaker database and check if any of them appears in the line extracted from the header. If it exists, at the first found I stop, and tag that name in the header and body as speaker. In case there is no information about the speaker in the header, first I split the body of the email into sentences using the '*sent\_tokenize*' function of *NLTK*. Then each sentence is split into words with *NLTK*'s '*word\_tokenize*' function, and then POS tagging (described at point 7.1) is applied. Next I apply a NER tagger (described at point 7.2) which will return an entity-tagged sentence in *nlTK.Tree* format. From that tree, I extract all the entities tagged as '*per*' or '*org*' and check if that entity exists in my speaker database. If it existed, I tagged it accordingly. This approach achieved a precision of **0.537**, recall of **0.258** and F1-Score of **0.34**.

Further improvements could be made on this algorithm given more time, by not only relying on my pre-build database, but checking on Wikipedia if a found entity by the NER tagger is a person or not.

## 6. Location Tagging

First I check whether there exists a line in the header starting with the word 'Place'. If it exists I tag the content after the word as location and also all its occurrences in the body. If no information is found in the header, I loop through my location database, which was constructed from all the location tags from the training data. If any of them appears in the body, I tag it accordingly. This algorithm achieved a precision of **0.91**, recall of **0.73** and F1-Score of **0.81**. Even though this algorithm performs really well on this sensitive data (location of seminar), it could be improved by training the NER tagger on more location-based text and then also using

some 3<sup>rd</sup> party library to check whether the identified entity is a valid location (GoogleMaps API).

## 7. Custom Algorithms

The reason why I decided to not use the in-built version of the following algorithms is because in my opinion nowadays sophisticated Machine Learning algorithms can outperform the old-fashioned NLP algorithms, and also the task was to build our algorithms as much as possible, to understand the core of NLP.

### 7.1. POS Tagger

As a POS tagger, I trained a Decision Tree Classifier using *scikit-learn*, which after learning from already annotated data, was able to POS tag new sentences. I used the ‘*treebank*’ corpus as data for my classifier, because it has a huge amount of data already tagged. Instead of just relying on features like the previous word (bigram) or the two previous words (trigram), after some research on the internet, I found a few relevant word features based on which my classifier should work, eg:

- is it the first word in the sentence?
- is it the last word in the sentence?
- is the first letter capital?
- is it all capitalized?
- prefixes
- suffixes
- previous word
- next word
- is it numeric?
- has hyphen?

(NLP-FOR-HACKERS,2017)

After gathering my training data and defining the features my classifier relies on, I had to prepare my data for training. For this reason I split the pre-tagged corpus into two arrays: the input, containing dictionaries of features about each word, in order as they appeared in the corpus; the output, which contained the expected POS tag for each word. Once the training was done, I was able to test it, since I split my corpus into 80%-20%, and the resulting accuracy was **94.5%**.

Because the training took about 8 hours, I had to save it to disk, and whenever an instance of the *POSTagger* class is created for use, the model loads up from disk. As input it accepts an array of words (word tokenized sentence) and outputs an array of tuples of the form (*word*, *tag*).

### 7.2. NER Tagger

My first attempt was to train a NER tagger using the in-built *NLTK* tools, but there were a few issues:

- the training corpus was too big, and the program was running out of memory, but from my perspective, in most of the cases more data means better accuracy
- even though the accuracy was 93%, I believed that I could achieve better results
- based on my POS tagger experience, I thought I could come up with better feature sets that better describe the data

For these reasons I decided to use Machine Learning again, to build a NER tagger. When choosing my model, I had to choose one which supports **Out-Of-Core Learning** (Scikit-learn.org, 2017). This is a learning process through which we keep the training data out of RAM. We rather load it in batches and do **Incremental Learning** with the classifier. I chose the Perceptron model, because it trains fast and also gives good results. As a dataset I used the ‘*Groningen Meaning Bank*’ corpus, which is already NER tagged. After pre-processing the data by separating the words from the resulting NER tags, I had to get documented on word features which could improve my classifier, and came up with:

- lemma of word
- is it capitalized?
- is it lowercase?
- is it a number?
- ending dot?
- next word and details about it
- previous word and details about it

and so on... (NLP-FOR-HACKERS, 2017)

I split my corpus again, into a 80-20 ratio and trained it. In the end the classifier was saved to disk, and whenever needed it can be loaded from file using ‘*pickle*’.

The accuracy achieved was much better than the first NER algorithm’s, resulting in **97%**.

## Reference:

1. NLP-FOR-HACKERS, 2017 – A blog about simple and effective Natural Language Processing. Available at: <http://nlpforhackers.io>
2. Scikit-learn.org, 2017 – scikit-learn: machine learning in Python – scikit-learn 0.19.1. Available at: <http://scikit-learn.org>